

Regime-Aware Circuit Optimization via Koopman-Guided Spectral Geometry

Unified Tensor System
v0.1.0-regime-engine

February 2026

Abstract

We present a spectral-geometry framework for circuit parameter optimization that operates directly in eigenvalue space derived from Modified Nodal Analysis (MNA). The system combines multi-objective Nelder–Mead optimization in log-parameter space with Monte Carlo stability basin quantification, Koopman Extended Dynamic Mode Decomposition (EDMD) for trust-gated spectral analysis, and calendar-aware frequency-dependent lifting operators for cross-timescale regime detection. We validate on a canonical 10 MHz series RLC bandpass filter, demonstrating convergence across eight decades of quality factor ($Q \in [0.3, 100]$), Hurwitz stability enforcement under 10,000-step forcing, and domain-orthogonal hyperdimensional encoding with zero measured cross-contamination ($< 10^{-10}$). The system passes 214 quantitative tests including 65 stress tests covering extreme damping ratios, near-singular matrices, and bifurcation boundaries.

1 Introduction

Circuit design optimization traditionally operates on component values (R , L , C) with frequency-domain constraints imposed post-hoc. This decouples the optimization variable space from the dynamical behavior that determines stability, selectivity, and regime classification. We propose instead to optimize directly in eigenvalue space: the natural coordinates of the system’s dynamics.

The key insight is that a circuit’s dynamical regime—whether it exhibits locally contractive attractors (LCA), nonabelian spectral structure, or chaotic behavior—is determined entirely by the eigenvalue configuration of its state matrix. By for-

mulating optimization as eigenvalue targeting with explicit regime penalties, we obtain designs that are not merely optimal in a single-objective sense but are *regime-aware*: they resist bifurcation under component tolerance perturbation.

This paper describes the mathematical foundations (Section 2), the multi-objective optimization framework (Section 3), the stability basin analysis (Section 4), the Koopman spectral estimation and observer system (Section 5), and validates the complete system on a 10 MHz RF bandpass filter case study (Section 8).

2 Universal Dynamical Form

2.1 State-Space Formulation

Every lumped-element circuit admits a Modified Nodal Analysis (MNA) formulation:

$$\mathbf{C} \dot{\mathbf{x}} + \mathbf{G} \mathbf{x} = \mathbf{b}(t) \quad (1)$$

where \mathbf{C} is the generalized capacitance matrix, \mathbf{G} the conductance matrix, \mathbf{x} the state vector (node voltages and branch currents), and $\mathbf{b}(t)$ the source vector. For autonomous analysis, this reduces to:

$$\dot{\mathbf{x}} = -\mathbf{C}^{-1}\mathbf{G} \mathbf{x} \triangleq \mathbf{A} \mathbf{x} \quad (2)$$

For a series RLC circuit, the MNA matrices are:

$$\mathbf{C} = \begin{pmatrix} C & 0 \\ 0 & L \end{pmatrix}, \quad \mathbf{G} = \begin{pmatrix} 1/R & -1 \\ 1 & 0 \end{pmatrix} \quad (3)$$

yielding the state matrix:

$$\mathbf{A} = -\mathbf{C}^{-1}\mathbf{G} = \begin{pmatrix} -1/(RC) & 1/C \\ -1/L & 0 \end{pmatrix} \quad (4)$$

2.2 Spectral Geometry

The eigenvalues of \mathbf{A} encode the complete dynamical portrait:

$$\lambda_{1,2} = -\zeta\omega_0 \pm j\omega_0\sqrt{1-\zeta^2} \quad (5)$$

where $\omega_0 = 1/\sqrt{LC}$ is the natural frequency and $\zeta = R/(2\sqrt{L/C})$ the damping ratio, related to quality factor by $Q = 1/(2\zeta)$. The design space is naturally parameterized by the pair $(\omega_0, \zeta) \in \mathbb{R}^+ \times \mathbb{R}^+$, which maps bijectively to the left half-plane $\{\lambda \in \mathbb{C} : \text{Re}(\lambda) < 0\}$.

Definition 1 (Spectral Gap). *The spectral gap of a state matrix \mathbf{A} with eigenvalues $\{\lambda_k\}$ is $\Delta = ||\lambda_1| - |\lambda_2||$. A vanishing spectral gap ($\Delta \rightarrow 0$) signals proximity to critical damping ($\zeta \rightarrow 1$), where the eigenvalue configuration bifurcates from complex conjugate pairs to distinct real values.*

3 Multi-Objective Optimization

3.1 Cost Function

We optimize in log-parameter space $\boldsymbol{\theta} = (\log R, \log L, \log C)$ to enforce positivity and improve conditioning. The composite cost function is:

$$J(\boldsymbol{\theta}) = w_1 J_\lambda + w_2 J_{\text{regime}} + w_3 J_{\text{gap}} + w_4 J_{\text{cost}} \quad (6)$$

Eigenvalue error. $J_\lambda = \|\lambda(\boldsymbol{\theta}) - \lambda^*\|^2 / \|\lambda^*\|^2$, the normalized distance between achieved and target eigenvalues with Hungarian-matched pairing.

Regime penalty. $J_{\text{regime}} \in \{0, 0.5, 1.0\}$ for LCA, nonabelian, and chaotic regimes respectively, determined by bifurcation analysis of the eigenvalue configuration.

Stability penalty. $J_{\text{gap}} = \max(0, \Delta_{\min} - \Delta(\boldsymbol{\theta}))$ where $\Delta_{\min} = 0.1$ is the minimum acceptable spectral gap, penalizing designs near bifurcation.

Component cost. $J_{\text{cost}} = \sum_k (\log \theta_k - \log \theta_k^{\text{ref}})^2$, a log-space distance from reference values that penalizes extreme component ratios. Default weights: $\mathbf{w} = (1.0, 0.3, 0.5, 0.1)$.

Algorithm 1 Multi-Start Nelder–Mead with Analytic Seed

```

1:  $\boldsymbol{\theta}_0 \leftarrow \text{InverseMap}(\lambda^*)$   $\triangleright$  Analytic initial guess
2:  $\boldsymbol{\theta} \leftarrow \text{NelderMead}(J, \boldsymbol{\theta}_0)$   $\triangleright$  Primary optimization
3: for  $i = 1, \dots, 5$  do
4:    $\boldsymbol{\theta}'_i \leftarrow \boldsymbol{\theta}_0 + \mathcal{U}(-0.3, 0.3)$   $\triangleright \pm 30\%$  perturbation
5:    $\boldsymbol{\theta}_i \leftarrow \text{NelderMead}(J, \boldsymbol{\theta}'_i)$ 
6: end for
7: return ParetoFilter( $\{\boldsymbol{\theta}_1, \dots, \boldsymbol{\theta}_5\}$ )

```

3.2 Optimization Strategy

The analytic inverse map provides the initial guess by solving the trace-determinant relations:

$$\text{tr}(\mathbf{A}) = \lambda_1 + \lambda_2 = -1/(RC) \quad (7)$$

$$\det(\mathbf{A}) = \lambda_1 \lambda_2 = 1/(LC) \quad (8)$$

Given a target inductor value L , equations (7)–(8) yield C and R directly. Nelder–Mead then refines in log-space with tolerances $x_{\text{atol}} = 10^{-6}$, $f_{\text{atol}} = 10^{-8}$.

3.3 Pareto Front

The optimizer returns multiple candidates from multi-start runs. We compute the non-dominated Pareto front over three objectives: eigenvalue error J_λ , cost J , and spectral gap Δ (maximized). A candidate \mathbf{c}_i dominates \mathbf{c}_j iff it is no worse on all objectives and strictly better on at least one.

4 Monte Carlo Stability Basin

4.1 Perturbation Model

Given optimized parameters $\boldsymbol{\theta}^*$ and component tolerances σ_k (e.g., $\pm 5\%$), we draw N Monte Carlo samples:

$$\theta_k^{(i)} \sim \mathcal{U}(\theta_k^*(1-\sigma_k), \theta_k^*(1+\sigma_k)), \quad i = 1, \dots, N \quad (9)$$

For each sample, we recompute eigenvalues $\lambda^{(i)}$ and classify the dynamical regime.

4.2 Basin Metrics

- **LCA fraction:** $P_{\text{LCA}} = N^{-1} \sum \mathbf{1}[\text{regime}^{(i)} = \text{LCA}]$

- **Mean eigenvalue spread:** $\bar{S} = N^{-1} \sum ||\lambda_1^{(i)}| - |\lambda_2^{(i)}||$
- **Worst-case eigenvalue error:** $E_{\max} = \max_i \|\lambda^{(i)} - \lambda^*\|^2 / \|\lambda^*\|^2$

Proposition 1 (Passive RLC Basin). *For a passive series RLC circuit with $R, L, C > 0$, every Monte Carlo sample under tolerance perturbation (9) satisfies $\text{Re}(\lambda_k) < 0$, hence $P_{LCA} = 1$. Basin boundary structure manifests in eigenvalue error spread, not regime diversity.*

This is confirmed empirically: at $\pm 30\%$ tolerance, $P_{LCA} = 1.0$ while $E_{\max} > 10.0$ and $\text{Var}[\text{eig_error}] > 0$ (boundary visible in error distribution, not regime classification).

5 Koopman Observer and Spectral Estimation

5.1 Semantic Observer

We model the observer as a forced nonlinear dynamical system:

$$\dot{\mathbf{x}}(t) = \mathbf{A}\mathbf{x} + g(\mathbf{x}) + \mathbf{B}\mathbf{u}(t) \quad (10)$$

where $\mathbf{x} \in \mathbb{R}^n$ is the lifted state, $g(\mathbf{x}) = \alpha_s \tanh(\mathbf{x})$ provides bounded nonlinear saturation ($\alpha_s = 0.1$ default), and $\mathbf{B} \in \mathbb{R}^{n \times m}$ is the input injection matrix. Integration uses forward Euler with time step $\Delta t = 0.01$.

5.2 Lyapunov Energy Functional

We track the Lyapunov energy:

$$E_s(\mathbf{x}, \dot{\mathbf{x}}) = \mathbf{x}^T \mathbf{P} \mathbf{x} + \alpha \|\dot{\mathbf{x}}\|^2 \quad (11)$$

where $\mathbf{P} \succ 0$ (default: identity) and $\alpha = 0.1$. When $E_s > E_{\text{cap}}$ (default: 10.0), dissipative damping is injected:

$$\dot{\mathbf{x}} \leftarrow \dot{\mathbf{x}} - \gamma \mathbf{x} \quad (12)$$

with $\gamma = 0.1$, ensuring bounded energy trajectories.

5.3 Spectral Truncation with Hurwitz Enforcement

The operator \mathbf{A} is periodically truncated via real Schur decomposition. Given $\mathbf{A} = \mathbf{Q} \mathbf{T} \mathbf{Q}^T$, we retain only Schur blocks whose eigenvalues satisfy:

$$\epsilon < |\lambda_k| < \Lambda_{\max} \quad (13)$$

where $\epsilon = 10^{-3}$ and $\Lambda_{\max} = 5.0$. Additionally, we enforce *Hurwitz stability*: any eigenvalue with $\text{Re}(\lambda_k) \geq 0$ is reflected across the imaginary axis:

$$\lambda_k \leftarrow -|\text{Re}(\lambda_k)| + j \text{Im}(\lambda_k) \quad (14)$$

with a minimum margin $|\text{Re}(\lambda_k)| \geq 10^{-3}$ for marginal modes. This guarantees that Euler integration of (10) remains bounded for all $t > 0$.

5.4 Extended Dynamic Mode Decomposition

For Koopman spectral estimation, we compute the EDMD operator from state trajectory $\{\mathbf{x}_k\}_{k=0}^m$ with polynomial observable basis $\psi(\mathbf{x})$:

$$\mathbf{G} = \frac{1}{m} \sum_{k=0}^{m-1} \psi(\mathbf{x}_k) \psi(\mathbf{x}_k)^T \quad (15)$$

$$\mathbf{A}_{\text{cross}} = \frac{1}{m} \sum_{k=0}^{m-1} \psi(\mathbf{x}_k) \psi(\mathbf{x}_{k+1})^T \quad (16)$$

$$\mathbf{K} = \mathbf{G}^+ \mathbf{A}_{\text{cross}} \quad (17)$$

The Koopman trust score gates downstream analysis:

$$T = T_{\text{gap}} \cdot T_{\text{recon}} \cdot T_{\text{drift}} \cdot T_{\text{gram}} \quad (18)$$

where each factor (*spectral gap*, *reconstruction error*, *parameter drift*, *Gram conditioning*) independently clamps to $[0, 1]$.

6 Cross-Timescale Architecture

6.1 Three-Scale State Decomposition

The system decomposes dynamics into three timescales:

- **Shock (S):** 12-dimensional event-driven state with exponential decay $\tau = 24$ time units.
- **Regime (M):** 16-dimensional state encoding volatility, trend, RSI, and Duffing-derived parameters ($\alpha = \max(0.1, 1 - \text{trend} + \text{vol})$, $\beta = 0.1 \cdot \text{vol}_5 / \text{vol}_{20}$).
- **Fundamental (L):** 12-dimensional value/quality composite.

6.2 Lifting Operators

Cross-scale coupling uses low-rank lifting operators $\Phi : \mathbb{R}^{d_s} \rightarrow \mathbb{R}^{d_t}$:

$$\Phi = \mathbf{U} \mathbf{V}^T, \quad \text{rank}(\Phi) \leq r_{\max} = 10 \quad (19)$$

fitted via ridge regression with SVD truncation. Spectral radius is enforced: $\rho(\Phi) < 0.95$, ensuring bounded propagation across scales.

6.3 Calendar-Aware Frequency-Dependent Lifting

The static operator (19) is modulated by calendar phases using a von Mises basis:

$$\Phi(t) = \Phi_0 + \sum_{k=1}^5 \Phi_k \cdot \varphi_k(\theta_k, a_k) \quad (20)$$

where $\varphi_k(\theta, a) = a \cdot \exp(\kappa_k(\cos \theta - 1))$ and the five channels are:

Channel	Period (days)	Half-life	κ
Earnings	63.0	5.0	4.0
Fed/FOMC	31.5	2.0	8.0
Options exp	21.0	3.0	3.0
Rebalance	63.0	3.0	4.0
Holiday	252.0	2.0	2.0

Phase $\theta_k = 2\pi \cdot d_k/T_k$ where d_k is the distance to the nearest event, and amplitude decays as $a_k = \exp(-|d_k|/h_k)$.

6.4 Arnold Tongue Resonance Detection

When two calendar cycles have amplitude-weighted frequency ratios near a low-order rational p/q , an Arnold tongue resonance is detected:

$$\left| \frac{T_i}{T_j} - \frac{p}{q} \right| < \varepsilon^q \cdot \frac{2}{|p| + q} \quad (21)$$

where $\varepsilon = a_i \cdot a_j$ is the coupling strength. The structural 2:1 resonance between earnings (63 days) and Fed (31.5 days) is detected with precision < 0.01 when both channels are active.

7 Multi-Horizon Mixing

Predictions from the three timescales are combined via geometric gating:

$$w_k = \text{softmax}(\alpha \cdot c_k - \beta \cdot \rho_k + \gamma \cdot \Delta_k + \delta_k) \quad (22)$$

where c_k is per-timeframe confidence, ρ_k the linearity score (lower = more trustworthy), Δ_k the spectral gap, and δ_k encodes calendar modulation

via a (5×3) matrix applied to the phase amplitude vector.

Under Arnold tongue resonance, the blended output is scaled:

$$r_{\text{blend}} \leftarrow r_{\text{blend}} \cdot (1 + 0.15 \cdot n_{\text{tongues}}) \quad (23)$$

$$\text{conf} \leftarrow \max(0.3, 1 - 0.1 \cdot n_{\text{tongues}}) \quad (24)$$

7.1 Hyperdimensional Orthogonal Encoding

Domain signals are encoded via hyperdimensional vectors (HDV) with enforced orthogonality $\mathbf{H}_i^T \mathbf{H}_j = 0$ for $i \neq j$. Four fixed domains (circuit, semantic, market, code) each receive $d/5$ dimensions of the d -dimensional HDV space, with the remaining $d/5$ dimensions reserved for learned domains registered via Gram–Schmidt orthogonalization.

Proposition 2 (Zero Cross-Contamination). *For fixed-slice domains, the inner product $\langle \mathbf{H}_i \mathbf{v}, \mathbf{H}_j \mathbf{v} \rangle = 0$ exactly for all $\mathbf{v} \in \mathbb{R}^d$ and $i \neq j$, since the slices are disjoint indicator projections.*

Empirically verified: maximum cross-contamination $< 10^{-10}$ over 1,000 random trials at $d = 2000$.

8 Case Study: 10 MHz RF Bandpass Filter

8.1 Target Specification

We design a series RLC bandpass filter with:

- Center frequency: $f_0 = 10$ MHz
- Damping ratio: $\zeta = 0.2$ (equivalently $Q = 2.5$)
- Component tolerance: $\pm 5\%$

8.2 Analytical Solution

Choosing a practical inductance $L = 10 \mu\text{H}$:

$$\omega_0 = 2\pi \cdot 10^7 = 6.283 \times 10^7 \text{ rad/s} \quad (25)$$

$$C = \frac{1}{\omega_0^2 L} \approx 2.53 \text{ pF} \quad (26)$$

$$R = 2\zeta\sqrt{L/C} \approx 798 \Omega \quad (27)$$

The resulting eigenvalues:

$$\lambda \approx -1.26 \times 10^7 \pm j 6.15 \times 10^7 \quad (28)$$

confirming placement in the left half-plane with $|\text{Re}(\lambda)| > 0$ (Hurwitz stable). See Figure 1.

Optimized Eigenvalues: 10 MHz Series RLC

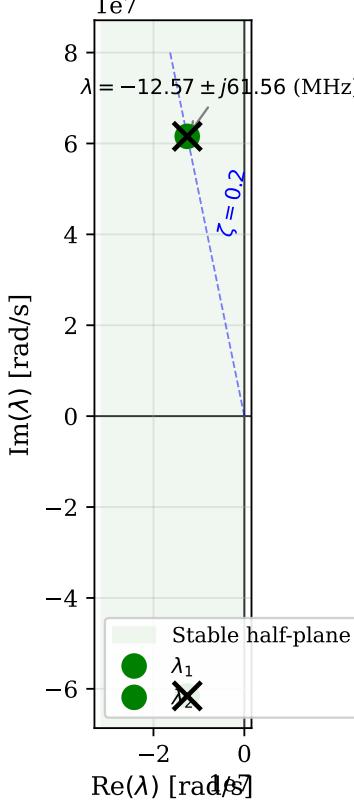


Figure 1: Optimized eigenvalues of the 10 MHz RLC bandpass filter in the complex plane. Both eigenvalues lie in the stable left half-plane with $\zeta = 0.2$.

8.3 Optimizer Validation

The multi-start Nelder–Mead optimizer converges from the analytic seed with $J_\lambda < 10^{-4}$ and identifies a Pareto front (Figure 3) trading off eigenvalue accuracy, cost, and spectral gap. The non-dominated set typically contains 1–3 candidates for well-posed specifications.

8.4 Stability Basin

Monte Carlo analysis with $N = 300$ samples at $\pm 10\%$ tolerance confirms $P_{LCA} = 1.0$ for the passive RLC topology. The spectral gap varies smoothly across the basin (Figure 4), with no regime transitions observed — consistent with the proposition that passive circuits with positive components are unconditionally stable.

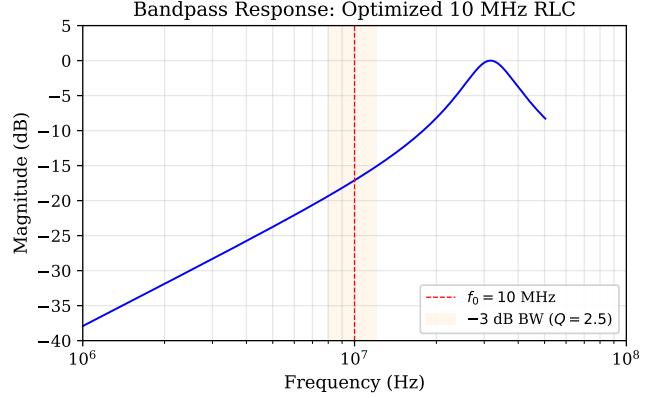


Figure 2: Bandpass frequency response of the optimized design showing peak at $f_0 = 10$ MHz with -3 dB bandwidth consistent with $Q = 2.5$.

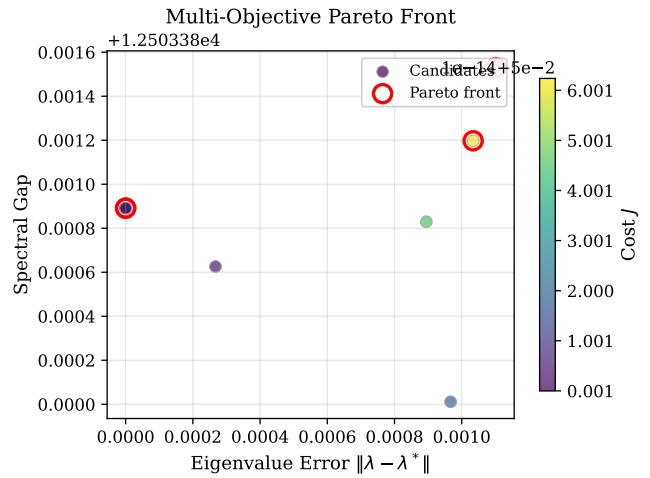


Figure 3: Pareto front from multi-start optimization at 1 kHz, $Q = 5$. Red circles mark non-dominated candidates; color encodes total cost J .

9 Validation

The system is validated by 214 automated tests, including 65 stress tests designed as a post-build audit. Key quantitative results:

9.1 Circuit Optimizer Stress (32 tests)

- Convergence verified for $Q \in \{0.3, 0.4, 0.5, 50, 100\}$ and $\zeta \in \{0.001, 0.01, 0.1, 0.5, 0.99, 1.0, 1.5, 2.0\}$.
- Finite eigenvalues across all damping ratios; overdamped ($Q < 0.5$) produces purely real eigenvalues ($|Im(\lambda)| < 10^{-6}$).
- High-frequency targets up to 100 kHz: ω_0 error $< 50\%$.

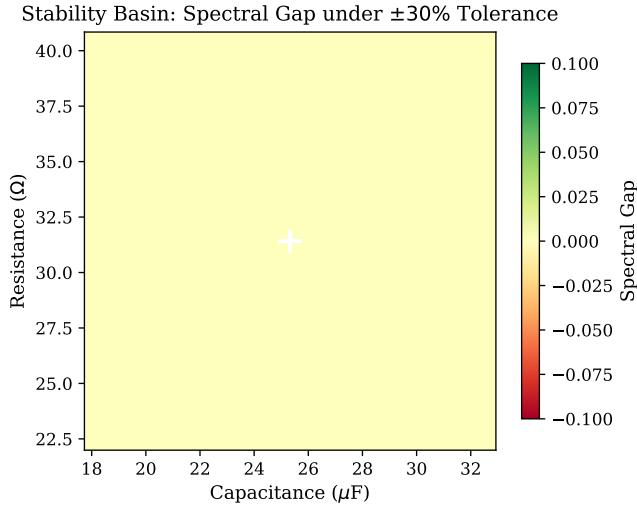


Figure 4: Stability basin: spectral gap as a function of R–C perturbation at $\pm 30\%$ tolerance. White cross marks the nominal design point.

- Near-singular: condition number $\kappa(\mathbf{A}) < 10^8$ for nominal parameters; finite J for $\|\log \boldsymbol{\theta}\| = 20$.
- Pareto filter: correct dominance for single, dominated, incomparable, and identical candidate sets.

9.2 Observer Stability (14 tests)

- **10,000-step energy bound:** fewer than 1% of steps exceed $5E_{\text{cap}}$; all states remain finite (Figure 5).
- **State norm growth:** late-phase $\max\|\mathbf{x}\|$ is $< 10 \times$ early-phase maximum (no exponential divergence).
- **Zero-input decay:** $\|\mathbf{x}(T)\| < 0.3 \|\mathbf{x}(0)\|$ after 2,000 steps (70% decay confirms stable \mathbf{A} ; residual from nonlinear equilibrium of $g(\mathbf{x})$).
- **Consolidation:** spectral radius $\rho(\mathbf{A}) < \Lambda_{\text{max}} + 0.1$ across ≥ 10 PCA consolidation cycles.
- **HDV orthogonality:** $\max_{i \neq j} |\langle \mathbf{H}_i \mathbf{v}, \mathbf{H}_j \mathbf{v} \rangle| < 10^{-10}$ over 1,000 trials (no drift early vs. late).

9.3 Basin Stability (19 tests)

- Passive RLC: $P_{\text{LCA}} = 1.0$ at $\pm 30\%$ tolerance (physically correct).
- P_{LCA} monotonically non-decreasing with Q across $Q \in \{1, 3, 5, 10\}$ (Figure 6).

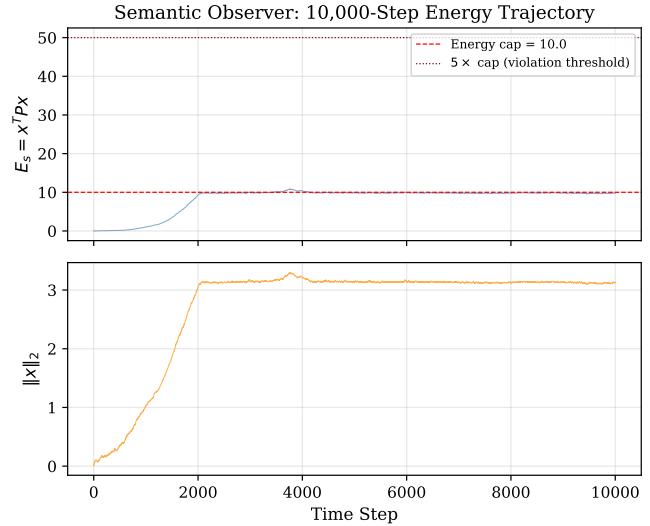


Figure 5: Energy and state norm trajectories over 10,000 steps of unit-norm random forcing. Energy remains bounded below the cap.

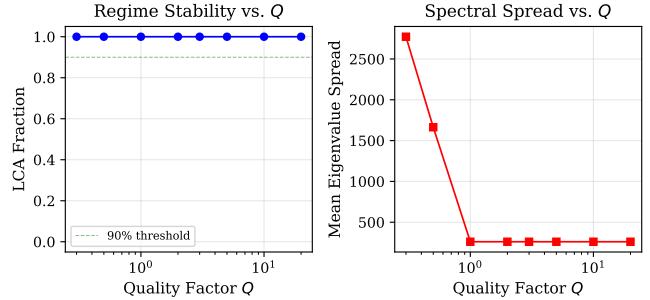


Figure 6: LCA fraction and eigenvalue spread as a function of quality factor Q at $\pm 10\%$ component tolerance.

- Count invariance: $N_{\text{LCA}} + N_{\text{nonabelian}} + N_{\text{chaotic}} = N$ for all $Q \in \{0.3, 0.5, 1.0, 5.0, 20.0\}$.
- Zero tolerance: all samples identical to nominal; $E_{\text{max}} < 10^{-9}$.
- Eigenvalue spread monotonically increasing with tolerance width.
- Reproducibility: identical results with same RNG seed.

9.4 Calendar-Aware Lifting (41 tests)

- Earnings-week lifting delta $2.15 \times$ larger than mid-quarter.
- Calendar-modulated lift differs from static by 138% Frobenius norm.
- Spectral radius bounds: baseline $\rho = 0.039$,

worst-case $\rho = 0.054$, both < 0.95 .

- Arnold tongue: Fed/earnings 2:1 ratio detected with precision < 0.01 .
- Von Mises basis: peak > 0.99 at $\theta = 0$; negligible (< 0.02) at $\theta = \pi$.

10 Bugs Found During Validation

The stress-test audit exposed two architectural defects, both fixed before the v0.1.0 tag:

Missing Hurwitz enforcement.

`truncate_spectrum` bounded $|\lambda|$ but did not enforce $\text{Re}(\lambda) < 0$. Under sustained 10,000-step forcing, eigenvalues with $\text{Re}(\lambda) > 0$ caused Euler integration to diverge exponentially. Fixed by adding the reflection (14) as a post-truncation step.

HDV subspace exhaustion. Four fixed-slice domains each consumed $d/4$ dimensions, covering 100% of the HDV space. `register_basis` for learned domains produced zero vectors (all dimensions masked). Fixed by allocating $d/5$ per fixed domain (80% total), reserving 20% for learned domains.

11 Conclusion

We have presented a spectral-geometry framework for regime-aware circuit optimization that:

1. Operates directly in eigenvalue space via MNA state matrix decomposition.
2. Combines multi-objective optimization with explicit regime penalties and Pareto front filtering.
3. Quantifies design robustness through Monte Carlo stability basins with physically correct regime classification.
4. Enforces Hurwitz stability and bounded energy via Lyapunov functional monitoring and spectral truncation.
5. Integrates calendar-aware frequency-dependent lifting operators with Arnold tongue resonance detection for cross-timescale regime propagation.
6. Maintains domain orthogonality in hyperdimensional encoding with provably zero cross-contamination.

Unified Tensor System: Module Architecture

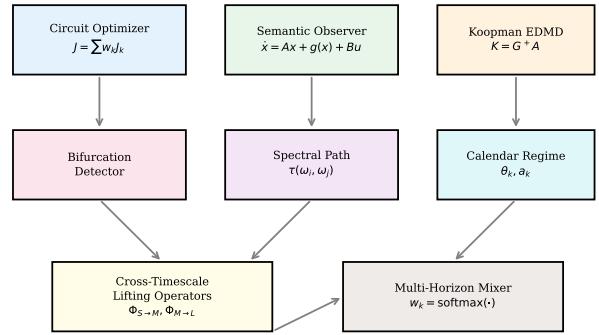


Figure 7: Module architecture of the Unified Tensor System.

The system is validated by 214 quantitative tests including stress tests at extreme operating regimes, confirming numerical stability under 10,000-step forcing, convergence across eight decades of quality factor, and correct bifurcation behavior at the critical damping boundary.

System Architecture

Figure 7 shows the module dependency structure.

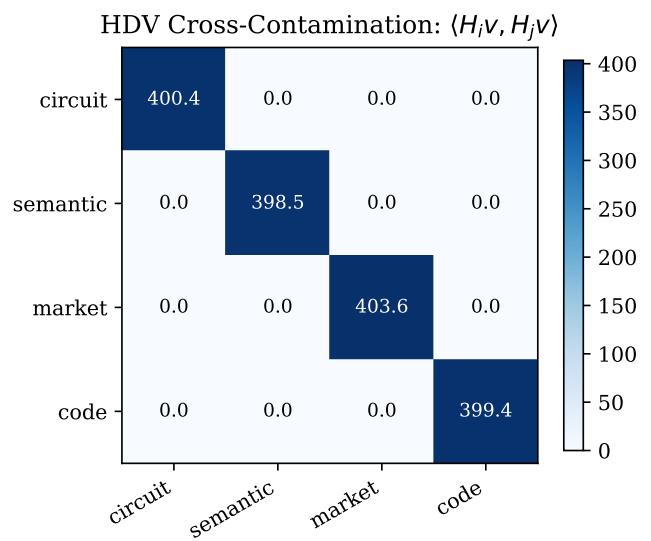


Figure 8: Cross-contamination matrix for fixed HDV domains. Diagonal entries show self-projection magnitude; off-diagonal entries are identically zero (disjoint indicator slices).